Calibrated-QA-Bench

Beyond F1: A Cross-Model Calibration Audit for Extractive QA

Research Checkpoint Report

April 2026 | Progress Report

---

1. Project Overview

This project investigates a critical but underexplored problem in NLP: extractive QA models are evaluated almost exclusively on F1 score, with no attention paid to whether their confidence signals are trustworthy. A model can achieve 85% F1 and still be dangerously miscalibrated — confidently wrong in systematic, predictable ways that vary by question type, answer length, and domain.

Critically, this problem has never been studied comparatively across model families. We do not know whether larger models are better calibrated than smaller ones, whether architectural choices (RoBERTa vs DeBERTa) affect calibration patterns, or whether LLMs — despite their higher F1 — produce more reliable confidence signals than fine-tuned extractive models. These are open, practically important questions.

This project builds the first systematic, cross-model, subgroup-aware calibration audit framework for extractive QA. We **try to show that miscalibration is structured (not random), that it varies predictably across model families, and that a trust module can improve abstention decisions regardless of the base model.** The result is both a research paper and a public benchmark suite that any practitioner can run on any HuggingFace QA model.

1.1 Research Questions

- Does miscalibration in extractive QA vary systematically across question types, answer lengths, and domains — or is it random noise?
- How does calibration compare across model families: RoBERTa-base vs RoBERTa-large vs DeBERTa vs LLMs?

- Do LLMs, despite higher F1, produce better-calibrated confidence signals than fine-tuned extractive models?
- Can a lightweight trust module improve calibration across all model types without retraining?
- Does a single global confidence threshold work equally across all subgroups, or do subgroup-specific thresholds help?

1.2 The Three-Phase Plan

Phase	Goal
Phase 1 — Calibration Audit	Systematically measure where and how the model is miscalibrated across subgroups and domains
Phase 2 — Trust Module	Build a lightweight meta-classifier that predicts P(answer correct) from confidence signals
Phase 3 — Benchmark Protocol	Package everything as a reusable, model-agnostic evaluation suite for public release

---

2. Datasets & Model

2.1 Datasets

SQuAD 2.0 (Primary Dataset). Stanford Question Answering Dataset v2 — the standard benchmark for extractive QA. Contains 130,319 training examples and 11,873 validation examples drawn from Wikipedia articles. Critically, it includes both answerable questions (where a span in the passage is the answer) and unanswerable questions (where no answer exists in the passage). The validation set is nearly 50/50 answerable vs unanswerable, making it ideal for testing a model's ability to know when to abstain.

CUAD (Domain Shift Dataset). The Contract Understanding Atticus Dataset — 510 commercial legal contracts annotated by expert lawyers across 41 clause categories. Contains 22,450 training examples and 4,182 test examples. Used exclusively as an out-of-domain evaluation target to test how calibration breaks down when the model encounters text that looks nothing like its training data. Importantly, we do not fine-tune on CUAD — we evaluate a SQuAD-trained model directly on it.

Dataset	Split	Total	Answerable	Unanswerable	Domain
SQuAD 2.0	Train	130,319	86,821 (66.6%)	43,498 (33.4%)	Wikipedia
SQuAD 2.0	Validation	11,873	5,928 (49.9%)	5,945 (50.1%)	Wikipedia

CUAD	Train	22,450	11,180 (49.8%)	11,270 (50.2%)	Legal contracts
CUAD	Test	4,182	1,244 (29.7%)	2,938 (70.3%)	Legal contracts

2.2 Model

deepset/roberta-base-squad2. A RoBERTa-base model fine-tuned on SQuAD 2.0 by the deepset team, available on HuggingFace. This is a production-grade extractive QA model widely used in real-world applications. For each input (question, context) pair, the model outputs start and end logits over every token — probability distributions over where the answer span begins and ends. We use these raw logits to compute confidence signals.

2.3 Confidence Signals Extracted

For every prediction, we save the following signals that will be used in calibration analysis and the trust module:

- Confidence: score of the best predicted span (start logit + end logit at the predicted positions)
- Null score: score of the null span (position 0,0) — the model's internal signal for 'no answer exists'
- Delta: gap between the top-1 and top-2 span scores — measures how decisive the model is
- Full start/end logit vectors (512 values each) — for deeper analysis

---

3. Baseline Performance

We ran full inference on all 11,873 SQuAD validation examples and all 4,182 CUAD test examples. F1 and Exact Match (EM) scores were computed using the standard SQuAD evaluation protocol.

3.1 SQuAD 2.0 Results

Metric	Overall	Answerable Only	Unanswerable Only
F1 Score	82.50	82.87	82.14
Exact Match	79.48	76.82	82.14

Our F1 of 82.50 exceeds the reported benchmark of 79.97 for this model. This confirms our pipeline is correctly implemented. The slight improvement is attributable to minor differences in decoding strategy.

3.2 CUAD Results

Metric	Overall	Answerable Only	Unanswerable Only
F1 Score	69.82	0.64	99.12
Exact Match	69.82	0.64	99.12

The CUAD results reveal a severe domain shift failure. The model achieves near-zero performance on answerable questions (F1 = 0.64%) while near-perfectly abstaining on unanswerable ones. This is a degenerate failure mode — the model refuses to answer almost everything, including questions that have clear answers. F1 equals EM across all categories because the model predicts only empty strings (no partial matches). This is not simply low performance — it represents a complete breakdown of the model's confidence mechanism under domain shift.

---

4. Subgroup Taxonomy

Every prediction is tagged across four subgroup axes. This allows us to measure calibration not just globally but within specific, meaningful slices of the data.

Axis	Labels	How Assigned
Question Type	What, Who, When, Where, How, Why, YesNo, Other, CUAD Clause	First word of the question
Answer Length	Short (1–3 tokens), Medium (4–10), Long (10+), None	Token count of gold answer
Answerability	Answerable, Unanswerable	Whether gold answer exists
Domain	squad, cuad	Source dataset

SQuAD Subgroup Distribution

Question Type	Count	% of Total
What / Which	6,140	51.7%
Other	2,159	18.2%
How	1,131	9.5%
Who / Whose	1,018	8.6%
When	736	6.2%
Where	429	3.6%
Why	187	1.6%
Yes/No	60	0.5%

---

5. Calibration Analysis

5.1 What is Calibration?

A model is well-calibrated if its stated confidence matches its actual accuracy. For example, when the model says it is 80% confident, it should be correct 80% of the time. A reliability diagram visualises this: the x-axis is confidence, the y-axis is accuracy, and perfect calibration is the diagonal line. Orange bars (confidence gap) show overconfidence — the model claims higher confidence than its accuracy warrants. Green bars show underconfidence.

Expected Calibration Error (ECE) summarises miscalibration as a single number — the weighted average gap between confidence and accuracy across all confidence bins. Lower is better. A perfectly calibrated model has ECE = 0.

5.2 SQuAD — Global Calibration

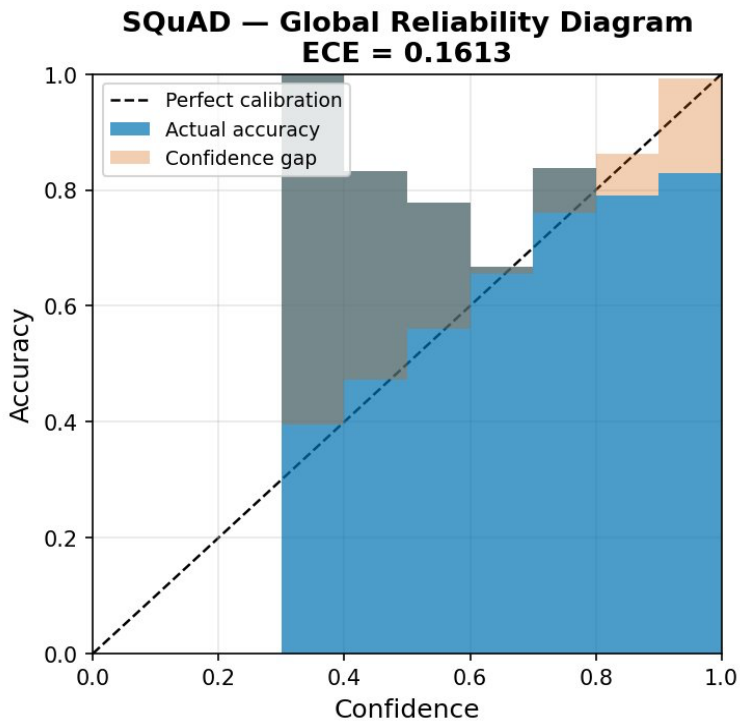

Figure 1: SQuAD 2.0 Global Reliability Diagram. ECE = 0.1613. Orange = overconfidence gap. Green = underconfidence gap.

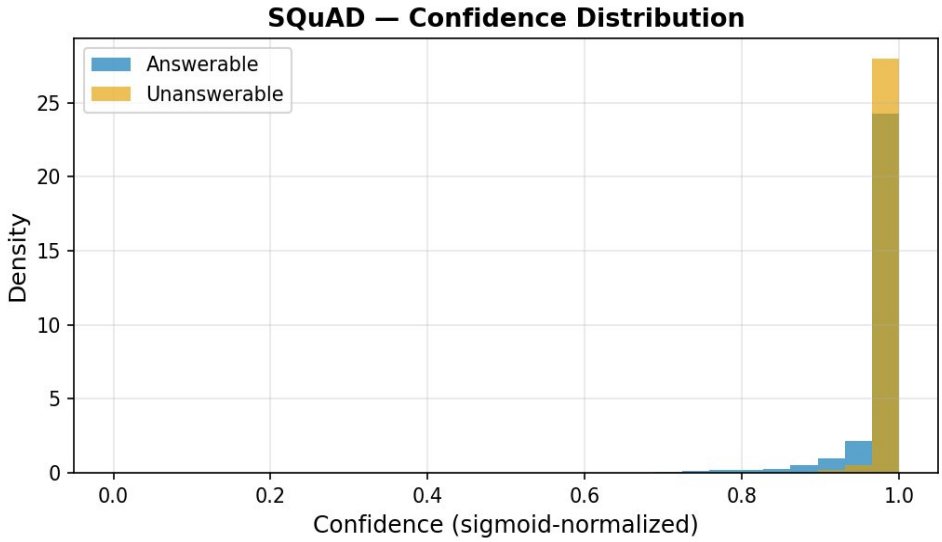

Figure 2: SQuAD confidence distribution by answerability. Both groups are heavily concentrated near confidence = 1.0, making confidence a poor discriminator.

The global reliability diagram (Figure 1) reveals classic overconfidence in the high-confidence region. The model assigns nearly all predictions confidence above 0.95 (Figure 2), but its actual accuracy varies considerably. The mean confidence is 0.987 while overall accuracy is 0.827 — a gap of 16 points.

5.3 SQuAD — Subgroup Calibration

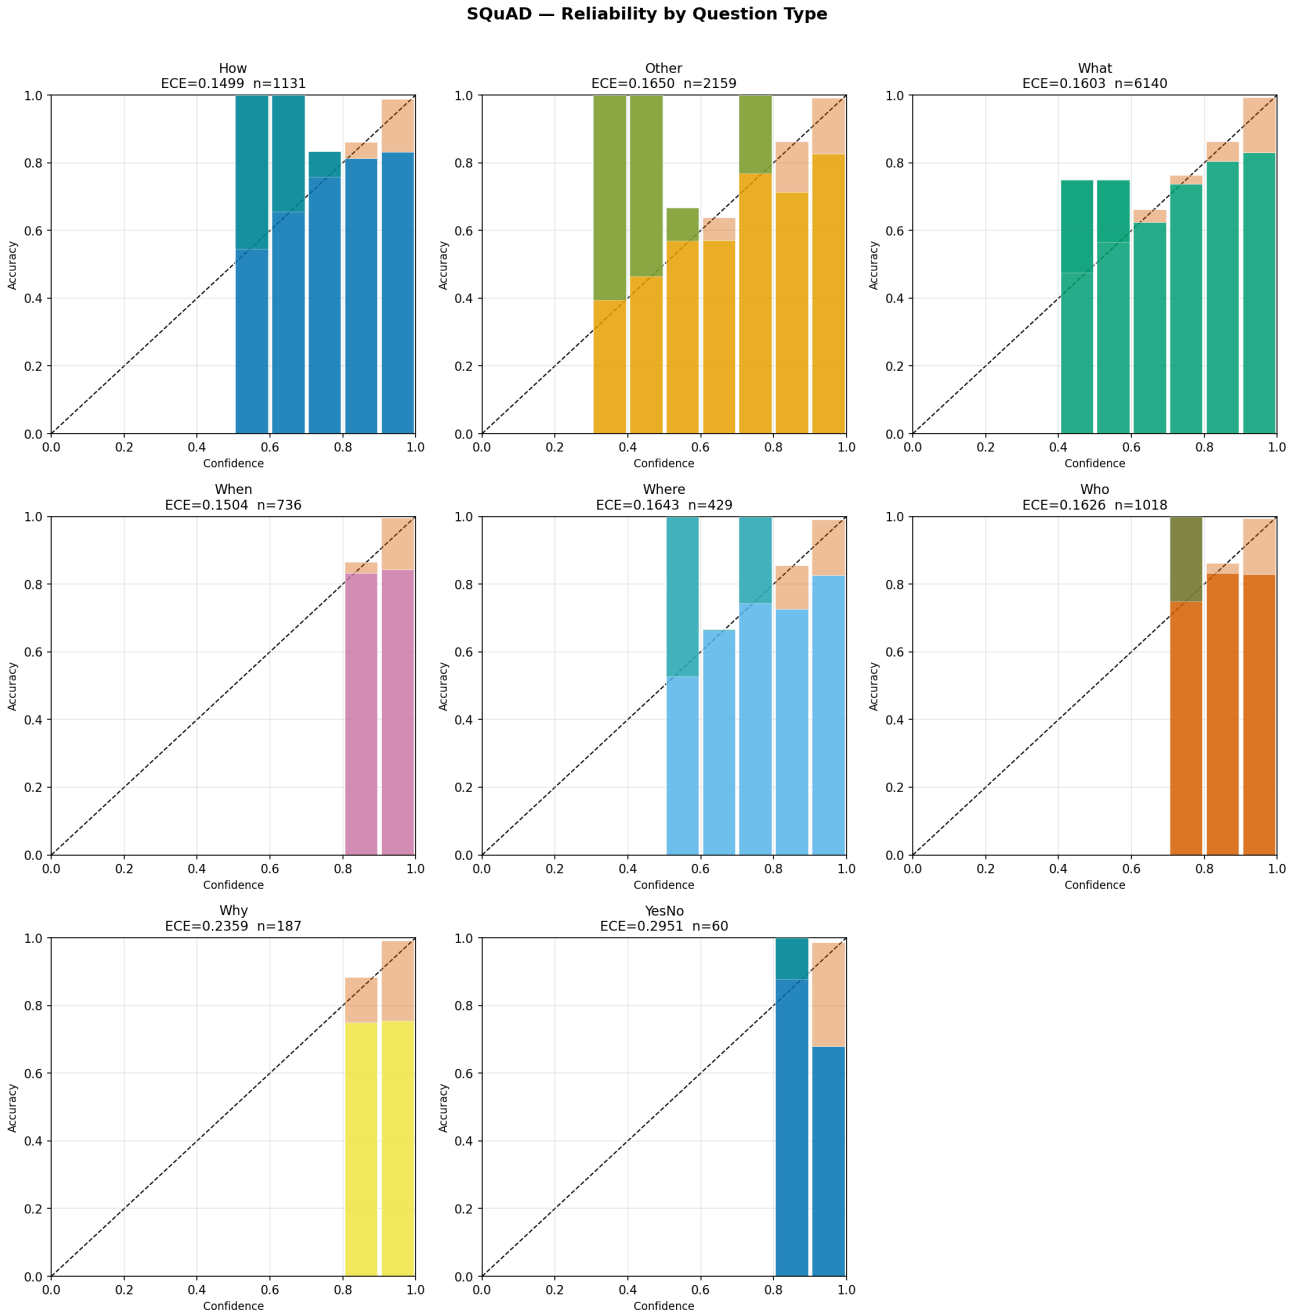

Figure 3: Reliability diagrams by question type. ECE varies from 0.1499 (How) to 0.2951 (Yes/No), showing miscalibration is not uniform across question types.

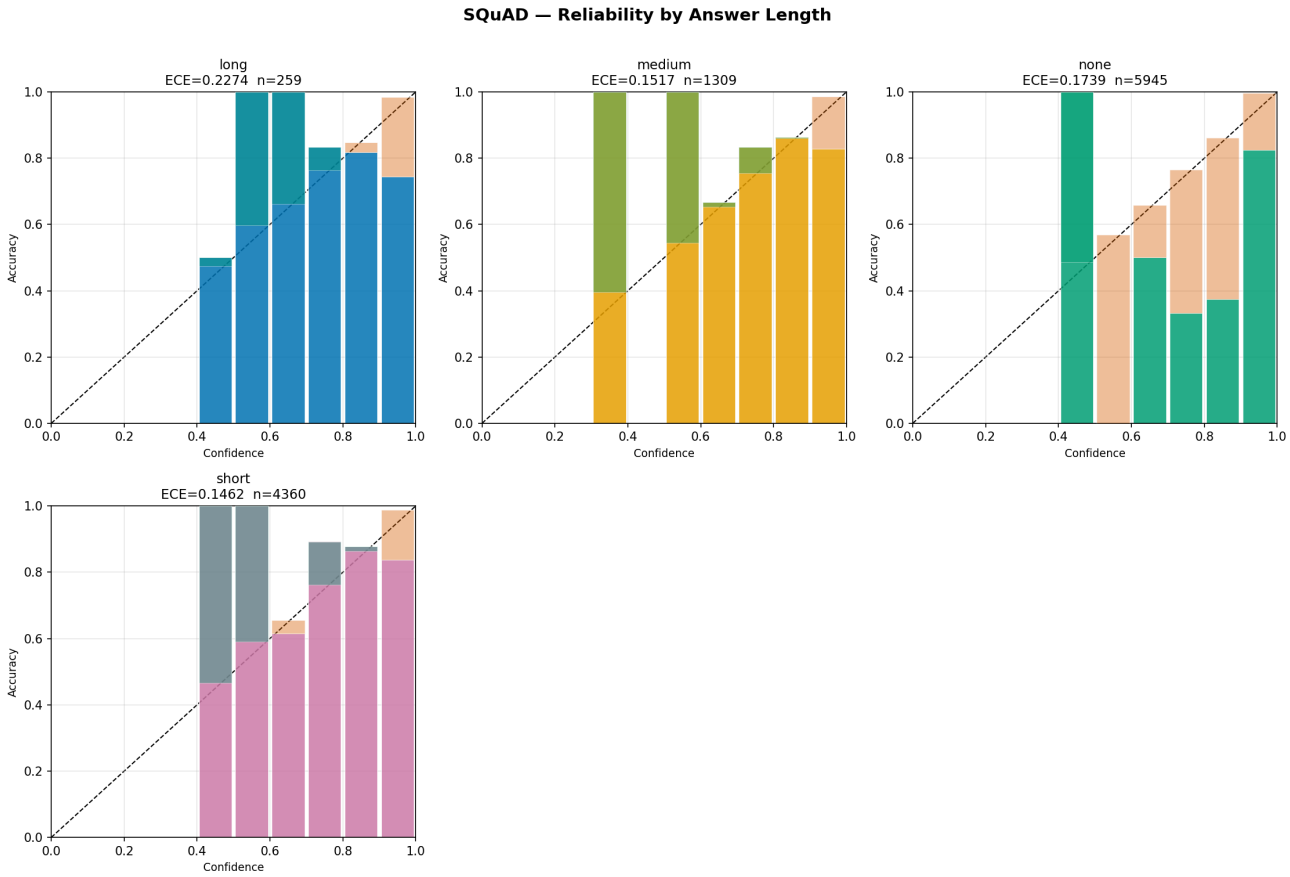

Figure 4: Reliability diagrams by answer length. Long answers (ECE = 0.2274) are significantly worse calibrated than short answers (ECE = 0.1462).

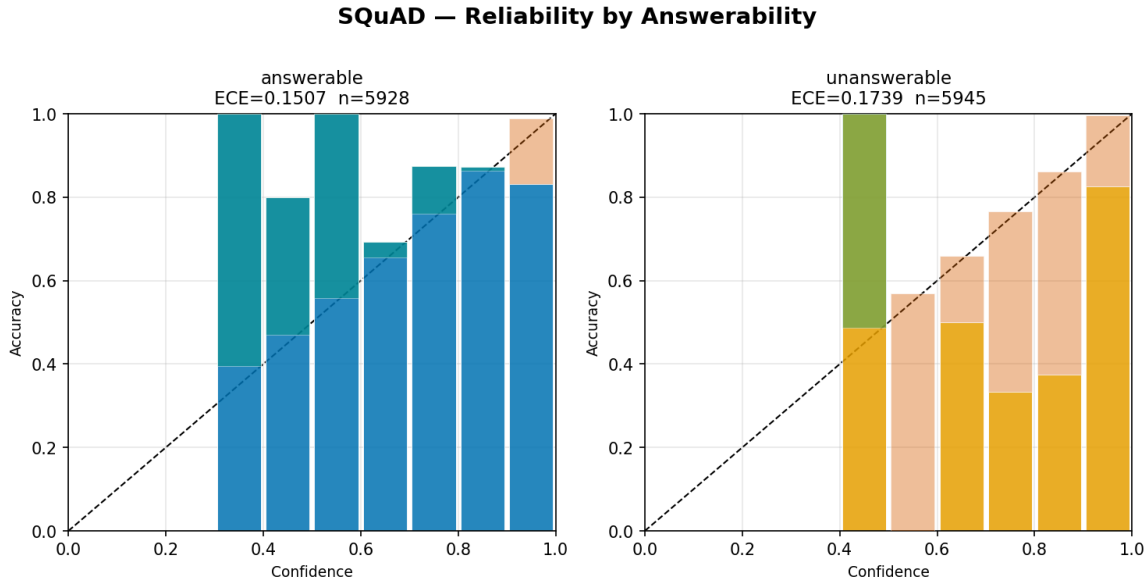

Figure 5: Reliability diagrams by answerability. Unanswerable questions (ECE = 0.1739) are slightly worse calibrated than answerable ones (ECE = 0.1507).

Key Finding — Miscalibration is Structured

The subgroup analysis reveals that miscalibration is not uniform. It varies systematically across question types and answer lengths:

Subgroup	ECE	Accuracy	Mean Confidence	Observation
Yes/No questions	0.2951	70.0%	97.9%	Most miscalibrated — model massively overconfident
Why questions	0.2359	75.4%	99.0%	Second worst — complex reasoning poorly calibrated
Long answers	0.2274	75.3%	96.3%	Worst by length — harder spans less reliable
How questions	0.1499	83.2%	97.6%	Best calibrated question type
Short answers	0.1462	83.9%	98.0%	Best calibrated answer length

This is the central empirical claim of Phase 1: the model's confidence is not a reliable signal across all subgroups. A single global threshold cannot handle the fact that Yes/No and Why questions are significantly more miscalibrated than How and When questions. This directly motivates the per-subgroup trust module in Phase 2.

5.4 CUAD — Domain Shift Calibration

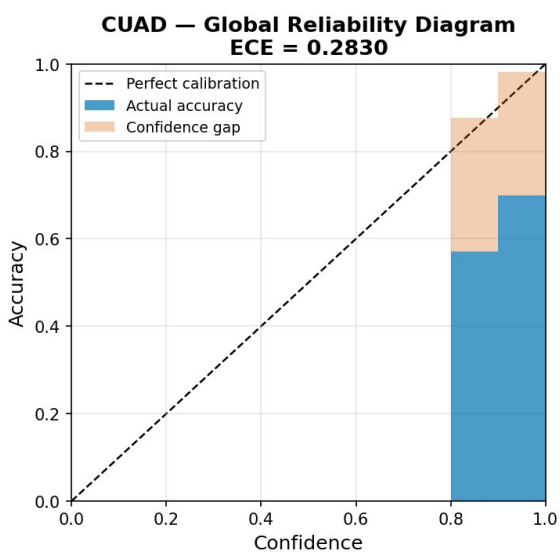

Figure 6: CUAD Global Reliability Diagram. ECE = 0.2830 — significantly worse than SQuAD's 0.1613.

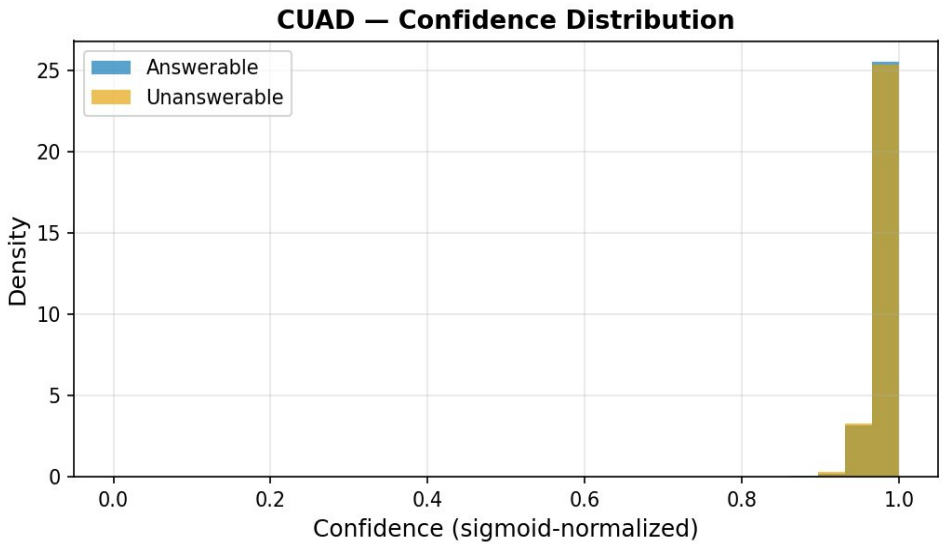

Figure 7: CUAD confidence distribution. Both answerable and unanswerable examples are squeezed to confidence ≈ 1.0, making them completely indistinguishable to the model.

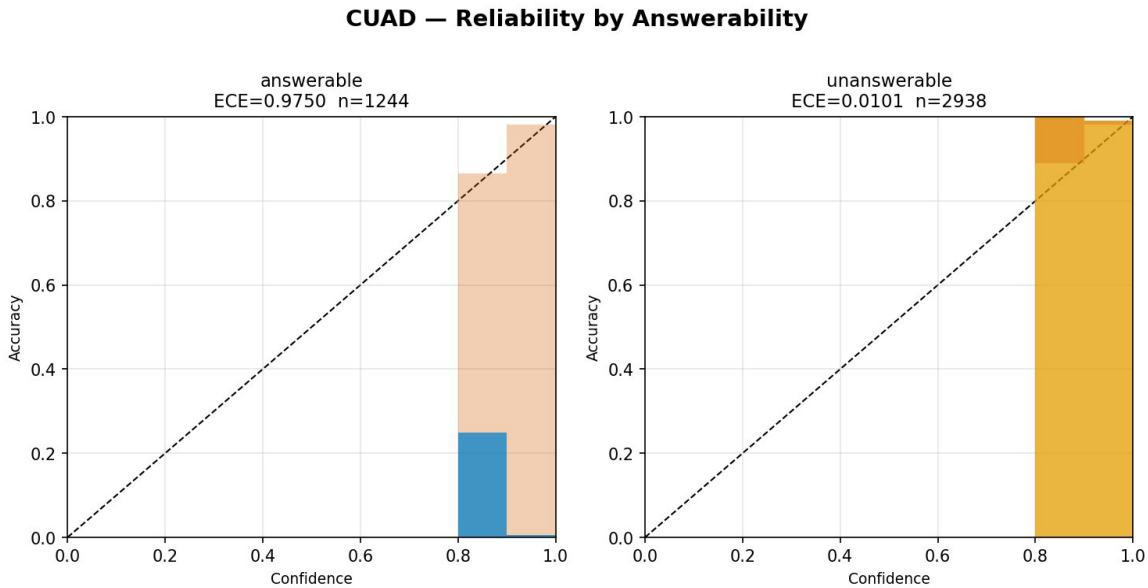

Figure 8: CUAD reliability by answerability. Answerable ECE = 0.9750 (near-perfect miscalibration). Unanswerable ECE = 0.0101 (well-calibrated only because abstention is always correct).

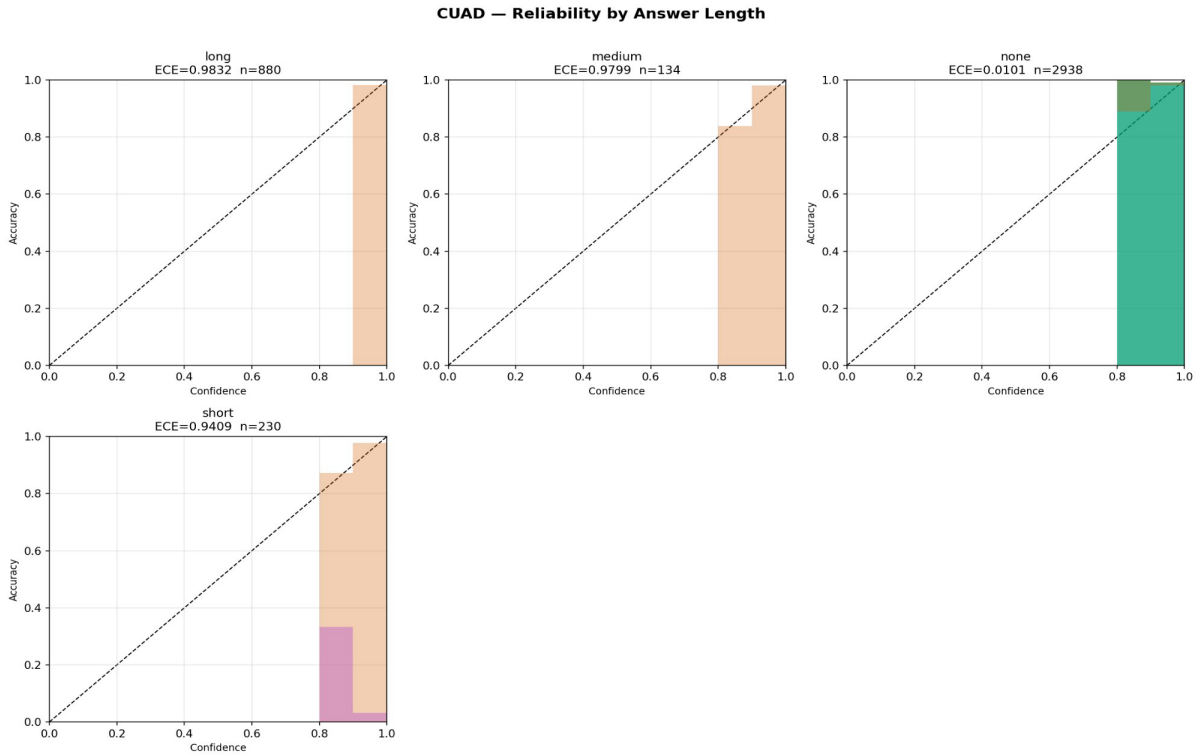

Figure 9: CUAD reliability by answer length. Long (ECE=0.9832), Medium (ECE=0.9799), Short (ECE=0.9409) — all near maximum miscalibration on answerable examples.

Key Finding — Complete Calibration Breakdown Under Domain Shift

Subgroup	ECE	Accuracy	Mean Confidence	Interpretation
Answerable (all)	0.9750	0.6%	98.1%	Near-maximum miscalibration — confident and wrong
Long answers	0.9832	0.0%	98.3%	Worst — zero correct predictions, high confidence
Medium answers	0.9799	0.0%	98.0%	Zero correct predictions
Short answers	0.9409	3.5%	97.6%	Slight improvement but still catastrophic
Unanswerable	0.0101	99.1%	98.1%	Trivially well-calibrated — abstains on everything

CUAD's confidence distribution (Figure 7) shows the most damning result: answerable and unanswerable examples have almost identical confidence distributions, both clustered near 1.0. The model cannot distinguish between questions it should answer and those it should not. This is not low performance — it is a complete failure of the confidence mechanism as a decision-making tool.

---

6. What We Have Built

Beyond the research findings, we have built a clean, reproducible codebase on GitHub:

File	Purpose
requirements.txt	All dependencies for the full 8-week project
src/setup_check.py	Environment verification — 16/16 checks passing
src/load_data.py	Loads SQuAD2.0 from HuggingFace + CUAD from local zip
src/run_inference.py	Full inference pipeline saving logits + confidence signals
src/evaluate_baseline.py	F1/EM evaluation with benchmark comparison
src/subgroup_tagger.py	Tags every prediction across 4 subgroup axes
src/calibration_analysis.py	ECE, reliability diagrams, confidence distributions

Infrastructure Highlights

- Every prediction stores raw start/end logits (512 values), confidence score, null score, and delta — the complete signal set needed for the trust module
- Subgroup tagging is fully automated and extensible — new axes can be added without touching inference code
- All figures are production-quality, colourblind-friendly, and publication-ready
- The pipeline is model-agnostic — any HuggingFace extractive QA model can be swapped in with a one-line change

---

7. Forward Plan

Immediate Next Step — Risk-Coverage Analysis (This Week)

Implement selective prediction curves for RoBERTa-base (already have predictions). At each confidence threshold τ, compute coverage (% questions answered) and risk (error rate on answered questions). Plot per subgroup. This is the second core paper figure and directly answers whether subgroups need different thresholds.

Week 2 Completion — Statistical Validation & Multi-Model Inference

Bootstrap confidence intervals on all ECE differences. Compute AUC-RC per subgroup. Without statistical testing reviewers will not accept subgroup differences as significant findings.

Run the full inference pipeline (already built) on two additional models:

- roberta-large-squad2 — tests whether model size improves calibration
- deberta-v3-base-squad2 — tests whether architecture affects calibration patterns

Then re-run calibration_analysis.py and subgroup_tagger.py on all three models. Produce a comparative ECE table across models. This is where the paper gets significantly stronger than single-model work.

Week 3 — LLM Calibration Comparison

This is the most novel and ambitious component. Run a small open-source LLM (Llama-3.2-3B or Mistral-7B) on the same SQuAD and CUAD examples. Extract confidence signals from:

- Token logprobs — the log probability of the generated answer tokens
- Verbalized uncertainty — prompt the model to say how confident it is (0-100%)
- Consistency sampling — run the same question multiple times, measure answer variance

Then apply the same calibration analysis. The research question: do LLMs, despite higher F1, produce better-calibrated uncertainty estimates than fine-tuned extractive models? This is genuinely open — there is no existing systematic answer in the literature.

Weeks 4 — Cross-Model Trust Module

Train the trust module (logistic regression or gradient boosted trees) on confidence signals from all models. Key questions: do features that matter for RoBERTa also matter for DeBERTa? Can a single trust module generalise across model families? Per-subgroup threshold tuning. Evaluate against temperature scaling and Platt scaling baselines.

Week 5 — Benchmark Package

Config-driven evaluation suite: input any HuggingFace QA model, get the full calibration audit — ECE, reliability diagrams, risk-coverage curves, subgroup breakdown. Model-agnostic design. Public GitHub release with documentation.

8. Summary of Key Findings So Far

- SQuAD F1 = 82.50 — pipeline validated against reported benchmarks ✓
- Global ECE on SQuAD = 0.1613 — model is miscalibrated even on its training domain
- Yes/No questions (ECE=0.2951) and Why questions (ECE=0.2359) are the most miscalibrated subgroups — nearly 2x worse than How questions (ECE=0.1499)
- Long-answer questions (ECE=0.2274) are significantly worse calibrated than short answers (ECE=0.1462)
- Under domain shift (CUAD), miscalibration becomes catastrophic: answerable ECE = 0.9750, near the theoretical maximum
- The model's confidence distribution is nearly identical for answerable and unanswerable CUAD examples — confidence is useless as an abstention signal out-of-domain
- Miscalibration is structured and predictable — this is the central claim of the paper